GPT-2 Based Meme Generator

Final Year Project (CTEC3451)
De Montfort University

Artem Bobrov (P2547788)

Jan 14 2021

Contents

Abstract

This project explores the possibilities of deep learning mechanisms to create and improve the ways of social interaction in the 21st century. The main idea of the project is to prove that modern trained neural networks can synthesise entertainment data based on the input, embedded in an application that is easy to use by every user.

Literature Review

Introduction

GPT-2 introduces wide possibilities when it comes to human-like text synthesis. It is developed by OpenAI in February 2019 but the release to the public was postponed due to the apprehension that it could potentially be used to conduct malicious and harmful activities. The successor to the first GPT family member has been trained on a larger dataset and parameter count with around ten-fold growth. The reason behind choosing this exact neural network to produce this project lies in the philosophy followed by its creators. The aim was to take a new approach to the purpose of the system. The past neural networks "are better characterized as 'narrow experts' rather than 'competent generalists'" (Radford et al. 2019) and that was the main restraining factor. GPT neural network family uses transformer architecture with attention mechanism which allows faster training due to significantly increased parralelism (Vaswani et al. 2017). There are several advantages of using transformer neural networks to process text when compared to convolutional and recurrent neural networks (CNN and RNN). Transformers deal with sentences within a text as a whole and not word by word. Self attention is another advantage of transformers since it allows to relate different positions of a single sequence (sentence) to calculate a representation of it (Vaswani et al. 2017). And third major difference is positional embeddings. This concept allows the usage of fixed weights that encode information about a specific token in a sequence. Convolutional neural networks are very good at processing images and detecting patterns in them. Recurrent neural networks are mainly used in evolutionary robotics to deal with vision (Harvey et al. 1994).

Transformer Neural Networks

Transformer neural networks continue to gain more popularity among the other artificial intelligence systems for solving 'Natural Language Processing' problems. Before transformers gated recurrent units with added attention mechanism were mainly used to process natural language but transformers based on solely that attention mechanism has proven that they're more effiecient in solving particular problems (Vaswani et al. 2017). A vast amount of them used encoder-decoder structure in which the encoder maps an input sequence of symbol representation to a sequence of continious representation. Using that continious representation the decoder generates an output sequence. Also it uses the already generated output as an additional input in order to learn and generate more accurate output.

Encoder and decoder stacks consist of six identical layers. Every encoder layer has two sub-layers and every decoder layer has three sublayers. The first two sublayers of the decoder are similar to the encoder sublayers. The first is a multi-head self-attention mechanism and the second is a fully connected feed-forward network. The third layer of the decoder works with the output of the encoder by performing multi-head attention over it.
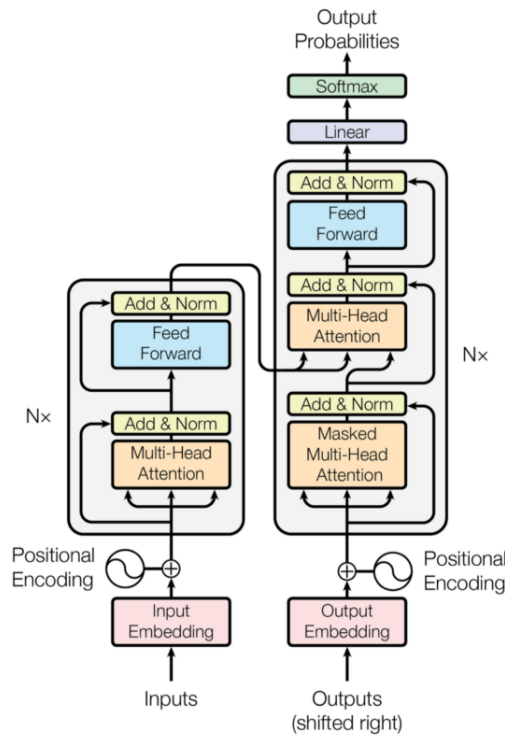

Figure 1: The Transformer - model architecture.

3

Since 2019 (release of GPT-2) a lot have changed in the field of transformers and as of November 2020 transformers are separated in different categories which different from each other in some aspects. Those categories are: Reformers, Linformers, Linear Transformers, Sinkhorn Transformers, Performers, Synthesizers, Sparse Transformers, and Longformers (Tay et al. 2020). Using this LRA benchmark the efficiency of different transformer models could be measured and this could give a better understanding how and in what areas they should be used. This testing suite is available in Python and is open-source so everyone could build upon it.

The suite developers pursued a goal of creating a suite for benchmarking having several important ideas in mind. Generality, Simplicity, Challenge, Long Inputs, Probing diverse aspects and accessibility of the suite. LRA Benchmark article depicts using a table of results for the comparison between different models. First test - Long 'ListOps' aims to test the parsing ability of the models and is a more complex version of the standard 'ListOps' task (Nangia & Bowman 2018). The best result is given by 'The Reformer'. Other conducted tests can be seen in the table provided in the research paper. The average score of 'Big Bird' is the highest, although it never showed best single result.

Convoluational Neural Networks

Convolutional neural networks were introduced long before transformers and have been mainly used to analyse visual imagery (Valueva et al. 2020). Also they've been proved to be able to undertake such tasks as recommender systems, natural language processing and financial time series. These networks are based on multilayer perceptron models but with implemented regularisation in order to avoid overfitting and that implies that multilayer perceptron neural networks are fully connected - each neuron of the previous layer is connected to every neuron in the next layer and that is the cause of overfitting. That problem has been solved with weight decay, which effectively is adding a penalty term to the cost function. This leads to weight shrinking in the backpropagation step, and trimming connectivity of a neural network which can be described as pruning - removing some unnecessary weights or neurons. Weight prunin is more popular than neuron because it's less likely to hurt network's performance and much easier to do.

The origin of the convolutional networks is the 'neocognitron' introduced by a Japanese computer scientist Kunihiko Fukushima in year 1980 (Fukushima 1980). It used two types of layers: convolutional layer and downsampling layer. The mechanism of interaction of these two types of layers is the following: Convolutional layer units whose receptive field cover previous layer. On the other hand downsampling layer units had receptive fields which covered the previous convolutional layers. The most popular method of training used for CNNs is backpropagation.

Convoluational neural networks perform unbelievably well in image recognition, for example - recogntion of handwritten characters testing using MNIST database showed an amazing result of 0.23% (Ciresan et al. 2012).

Recurrent Neural Networks

Recurrent neural networks have been derived from feedforward neural networks. They have internal state memory that allows them to work with long sequences of data. Everything that is sequential like connected handwriting or speech can be easily converted to digital format using RNNs (Abiodun et al. 2018). This type of networks is referred to as infinite impulse response class which can be applied not only to neural networks but to many other systems that produce an output based on an input which depends on time variable. There are different architecture variants of RNNs - fully recurrent, Elman and Jordan networks, Recursive and one that marked a big growth in recurrent neural networks - 'Long short-term memory' (Hochreiter & Schmidhuber 1997). LTSM networks showed a great performance in speech recognition, machine translation and other important areas (Sutskever et al. 2014).

LTSM also is a type of a recurrent neural network that happened to avoid 'vanishing gradient problem'. The problem affects the training process and weight adjustment in particular. The standard training process of a neural network is performed in a number of steps: the error is updated accordingly every iteration and the weight are changed regarding the error value. The problem is encountered when the gradient is so small that it prevents the weights from changing. The opposite of the 'vanishing gradient problem' is 'exploding gradient problem' - the change in weights is too big (Bohra 2021).

Applications of neural networks in entertainment

Introduction

Since artificial intelligence systems, transformer neural networks in particular, are good in synthesising text that is nearly indistinguishable from human phrases - that allows to make a presumtption: are modern neural networks developed and trained enough to be used to generate entertainment short texts based on already created so known 'memes'? It's important to understand the modern culture of internet and the heart of it - thematic pictures with short jokes written on them to undertake a project directly related to it. Memes can be found everywhere on the Internet nowadays and they constantly change and evolve. Their primary function is to entertain the community in a harmless way yet our society is very diverese and has different interests - that explains the existence of memes which can be used to "further political ideals,

maginfy echo chambers and antagonise minorities" since it became a very popular communcation method (Peirson & Tolunay 2018).

Tesseract

'Tesseract' is an optical character recognition tool which can be both electronic or mechanical and is used to convert many types of text (handwritten, typed, printed etc.) to digital (Kay 2007). A plenty of systems use this tool to scan all kind of documents to store it digitally. Modern OCR is trained enough to work with different fonts, languages and even image recognition whereas early versions had to be fed with every character separetely.

PyTesseract is a library written for Python which allows the usage of OCR functions directly in a program. This is vital for this project because the main source of data exists as images rather than plain text strings.

Pillow

'Pillow' is a library for Python that explores vast possibilities when it comes to working with images. Initially the project was called PIL (Python Imaging Library) but then it was discontinued in 2011. Pillow is a successor of the first library and is still being updated. The latest stable version was released in April 2021. The author of the fork (a copy of a repository that is being managed by someone else rather than the creator of the original repository) has set several goals to follow when maintaining the product - continuous integration testing, publicised development activity and regular releases to Python Package Index (Clark & Contributors 2022). It is capable of manipulating images in a variety of ways but one of the most important features of this library that will be used in this project is: "Adding text to an image".

PyTorch or TensorFlow?

PyTorch and Tensorflow are both tools for working with neural networks and offer different advantages and disadvantages. They both allow utilising your GPU's power in order to train NNs on big datasets. Any of these frameworks can be used to create an architecture of a neural network but they have some differences. GPU processing in PyTorch is implemented using CUDA system created and maintained by NVIDIA.

Conclusion

Summing up everything that has been stated before, nowadays the interest towards neural networks and other artificial intelligence systems grows with exponential speed. Literature available online that is crucial to familiarise yourself with when working

on such a project is extensive and allows more in-deep research to be carried out. Having that in mind, writings about AI are more inclined towards scientific purposes rather than entertainment. There are many types of neural networks available open-source: Recurrent Neural Networks, Convolutional Neural Networks, Perceptrons, Transformers etc. Transformers are the best choice when it comes to entertainment and text generation in particular, since they use attention mechanism - it greatly reduces training times. GPT-2 has a major training base of 1.5 billion parameters compared to BERT's two versions: 110 million and 340 million parameters (as of 2019). It makes GPT-2 a winner in more accurate text synthesis.

System Design

System UI (frontend)

The main key of the system UI when it aims to hold a 'dialogue' with the user is simplicity. The system can include a lot of parameters, functions, features which will be implemented in an outstanding way but if it fails to correctly introduce these features to the user and they become convoluted - it's a problem.

The ladning page features a simple approach to the design. In the beginning we have a context menu which will contain drop-down boxes to control the behaviour of the program which will be used to tailor the experience that suits the end user. It's planned the menu to contain settings to control the environment - change the appearence of the app, switch to another section (meme pattern), reload the app (in case of an emergency).

Next section of the interface is 'Branding'. The app shouldn't only be practical but also has a decent appearence. The logo should match to overall design of the program and to complement it.

'Directions to use' shows the user how to get the most out of the app and have a good experience. It will be shown only once on startup not to bother the user too much.

'Select image' section will allow the user to select the beloved meme pattern on which will be based the generation of the future memes shown. If the user doesn't have a preferred pattern or just wants to watch all kinds of memes - random generation will be available as an option (it will be set automatically, if no input collected).

The philosophy of the whole user interface which will be used in this project is to make the experience continious. Loading different pages manually is in the past, the modern idea of a scroller app is to make the usage seamless. Although the devices nowadays are very powerful and capable of handling massive amounts of

Figure 1: Main page design

calculations, it's still important to implement a good optimisation to make the usage of the app pleasant to the highest number of people. To balance somewhere in between experience and optimisation - 'lazy loading' is the answer. The main idea is to move away from separate pages and implement everything combined in one long page which will be loaded consequentially and automatically when the user reaches the end of the 'page'.

Also as can be seen on the diagram above, every picture (meme) should have several buttons to express opinion. This can be achieved by using typical 'like-dislike' system. To make even easier for the user to give positive feedback a double click system on the image can be used.
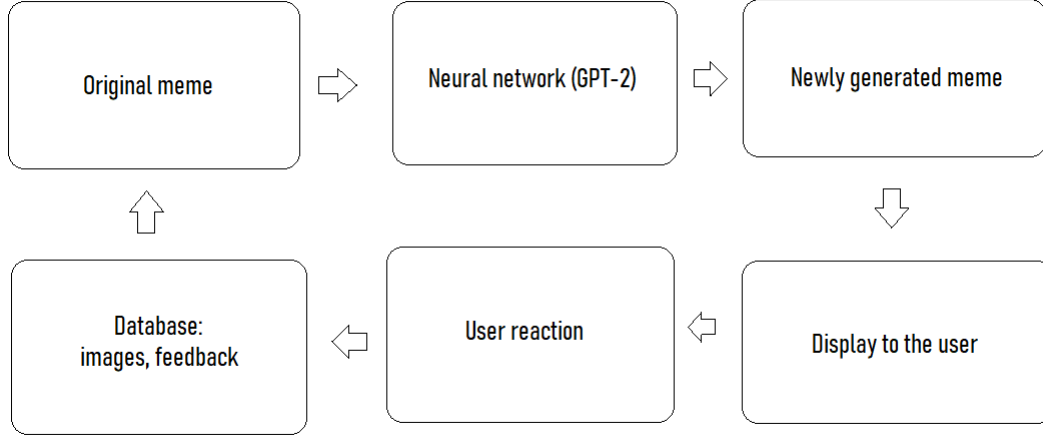
Interaction between frontend and backend

Figure 2: The system design flow chart

As can be seen on this figure the process involved in generating and showing the images is continous. First comes the database which will contain images and user reactions to each image. Reactions will be stored as a mass, not individual data. This will allow to understand better what the users like the most and what they don't. The program will take the initial image, separate the text from it, remember what the image was, feed the text to the neural network and combine the output from the neural network and the existing image into a freshly generated meme. After that comes the process of displaying the image. Then the user, if they wish, can leave feedback. The loop repeats.

The main bridge between backend and frontend in this project is the Python library called 'Eel' (Knott 2021). It is similar to JavaScript's 'Electron' and allows to code the frontend with a stack of HTML, CSS and JavaScript having Python in the backend. There is another library which is heavier and therefore has a higher entry level - 'cefpython'. 'Eel' makes possible linking buttons created using HTML/CSS to Python-implemented functions. This methodology is called 'Function Exposure' (Knott 2021). Functions from Python could be exposed to JavaScript environment and vice versa.

Backend design

Figure 3: Backend design

The backend of the project is going to be implemented using Python programming language due to the availability of different frameworks and libraries. In the backend the process of the scanning the text from the image will be undertaken by a library called 'PyTesseract' which allows the optical character recognition to be used in an optimised way. Next step of the backend process is to push the string to the neural network allowing it to use its massive database to synthesise similar text. Afterwards comes the process of attaching the generated output to an image. It is done using a tool called 'Pillow' and it's aim is to work with all kind of images. The string will be styled to fit the pattern appropriately.

Functional Requirements

The function requirement report will join the early described diagrams in a whole idea. Figure 4 shows the relationship between the user and the developer and how this interaction will enhance the experience of the user.


Figure 4: Use Case Diagram

The user has a way to communicate with the developer - hold a dialog in a non-verbal manner and explain what they like. This leads the system to be implemented in a 'prosumer' way. The user not only consumes the content but also influences it. The diagram explains different elements of the system and how they work together. Although the user has only one way to interact with the program, the program collects non-sensitive data about the user and decides which path to take. The 'use case diagram' is not final and may be changed throughout the development.

Overview	
Title	Scroll Example
Description	The user logs in to the app and scrolls down in order to see images
Actors	User
Pre-Conditions	First time login, no preferred image pattern selected
Basic Flow	
Step 1	Run the script to check the boolean variable for first login
Step 2	Show the welcome/landing page with short, simple instructions
Step 3 (1/2)	As the user scrolls down - generate new images, when the user reaches last 10 items (pre-generation to avoid lags)
Step 3 (2/2)	If the user doesn't reach to last 10 images - idle
Step 4	Unload the images that are far behind to decrease the load on the system
Step 5	As the user likes or dislikes some images - store the reaction linked to a certain image as statistical data
Post Condition	
Condition 1	The app saved the data about user login
Condition 2	Adjusted settings saved - will be used on next login

Table 1: Use Case Specification

The table above illustrates

Indicative Test Plan

Implementation Report

The approach chosen to develop this project is FDD - Feature-driven development. Features are implemented first and then the unit tests will be written to prove that these features work correctly. The whole idea of the project is to prove that neural networks can create memes that are interesting to people.

The current state of the development has the frontend and backend partly implemented. As of now the frontend features the ability to test the backend efficiently without writing unit tests. It visually shows what works and what doesn't. Content loading of the image stream when the user scrolls down is implemented using a 'Load More' button powered by jQuery. The app has a branding and instruction screen but since the backend is only partly implemented - the app doesn't store information about users and wether they're using it for the first time. In the backend some functions that process images are fully implemented. For example, functions that scan the text and add it to the images work but they have to be enhanced. Reading the text from the image sometimes collects non-existent characters and some kind of a filter should be implemented that will throw away useless characters. On the other hand the code that adds the text to the image works perfect.

For the frontend it's planned to implement 'lazy loading' using either React or jQuery and load the images from the database endlessly. For the backend the main

target is to adapt GPT-2 code and use it to process text.

Bibliography

Abiodun, O. I., Jantan, A., Omolara, A. E., Dada, K. V., Mohamed, N. A. & Arshad, H. (2018), 'State-of-the-art in artificial neural network applications'.
Accessed Jan 11 2022 [Online].

Bohra, Y. (2021), 'The challenge of vanishing/exploding gradients in deep neural networks'.
Accessed Jan 12 2022 [Online].

Ciresan, D., Meier, U. & Schmidhuber, J. (2012), 'Multi-column deep neural networks for image classification'.
Accessed Jan 11 2022 [Online].

Clark, A. & Contributors (2022), 'From pillow documentation'.
Accessed Jan 8 2022 [Online].

Fukushima, K. (1980), 'Neocognitron: A self-organizing neural network model for a mechanism of pattern recognition unaffected by shift in position'.
Accessed Jan 11 2022 [Online].

Harvey, I., Husbands, P. & Cliff, D. (1994), Seeing the light: Artificial evolution, real vision, pp. 392–401.
Accessed Jan 7 2022 [Online].

Hochreiter, S. & Schmidhuber, J. (1997), 'Long Short-Term Memory', *Neural Computation* **9**(8), 1735–1780.
URL: *https://doi.org/10.1162/neco.1997.9.8.1735*

Kay, A. (2007), 'Tesseract: an open-source optical character recognition engine'.
Accessed Jan 8 2022 [Online].

Knott, C. (2021), 'Eel github'.
Accessed Jan 10 2022 [Online].

Nangia, N. & Bowman, S. (2018), 'Listops: A diagnostic dataset for latent tree learning'.
Accessed Jan 7 2022 [Online].

Peirson, A. L. & Tolunay, E. M. (2018), 'Dank learning: Generating memes using deep neural networks'.
Accessed Jan 7 2022 [Online].

Radford, A., Wu, J., Child, R., Luan, D., Amodei, D. & Sutskever, I. (2019), 'Language models are unsupervised multitask learners'.
Accessed Jan 6 2022 [Online].

Sutskever, I., Vinyals, O. & Le, Q. V. (2014), 'Sequence to sequence learning with neural networks'.
Accessed Jan 12 2022 [Online].

Tay, Y., Denghani, M., Abnar, S., Shen, Y., Bahri, D., Pham, P., Rao, J., Yang, L., Ruder, S. & Metzler, D. (2020), 'Long range arena: A benchmark for efficient transformers'.
Accessed Jan 7 2022 [Online].

Valueva, M., Nagornov, N., Lyakhov, P., Valuev, G. & Chervyakov, N. (2020), 'Application of the residue number system to reduce hardware costs of the convolutional neural network implementation', *Mathematics and Computers in Simulation* **177**, 232–243.
URL: *https://www.sciencedirect.com/science/article/pii/S0378475420301580*

Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A., Kaiser, L. & Polosukhin, I. (2017), 'Attention is all you need'.
Accessed Jan 7 2022 [Online].